

# Juan PINEDA-JARAMILLO, PhD

AI Engineer | LLM Engineer – Agentic Systems, RAG & Production GenAI

[juandpineda@gmail.com](mailto:juandpineda@gmail.com) 

[jdpinedaj](#) 

[jdpinedaj](#) 

## AREAS OF EXPERTISE

- **Agentic AI & LLM Engineering:**
  - Multi-Agent Orchestration: Production multi-agent systems and multi-step agentic AI workflows with LangGraph, LangChain, the OpenAI Agents SDK and Google ADK (Agent Development Kit), including agent-to-agent (A2A) protocols for cross-domain evidence exchange over shared state. Supervisor and router patterns, tool calling, agent memory, and self-correcting agent loops.
  - Prompt Engineering & Structured Output: Prompt-as-configuration (versioned YAML templates), few-shot / zero-shot design, model grounding, strict Pydantic schema enforcement, model routing and fallback. Hands-on with Azure OpenAI (GPT-5.x), OpenAI, AWS Bedrock (Claude, Titan, Nova Pro) and HuggingFace models.
  - RAG & Retrieval: GraphRAG over Neo4j knowledge graphs, vector search with pgvector and FAISS, semantic example selection, multimodal ingestion (PPTX, PDF, DOCX, XLSX and slide-to-image vision), and enforced source citation for evidence traceability.
  - LLM Safety & Guardrails: Defence-in-depth guardrail frameworks: prompt-injection detection (regex and semantic), PII/PHI redaction (Presidio, spaCy), data-exfiltration and system-prompt-leakage blocking, schema and citation validation, per-run cost and token ceilings, model-version pinning, and immutable hashed audit trails.
  - Evaluation & Observability: Ground-truth eval harnesses, deterministic oracles combined with LLM-as-judge, adversarial red-team corpora, offline regression suites driven by mock LLMs, tracing with Arize Phoenix, OpenTelemetry and LangSmith, and per-run token, latency and cost tracking.
  - Text2SQL & Neuro-Symbolic Systems: Schema-agnostic SQL agents that discover schemas at runtime and recover from their own errors, combining LLM reasoning with deterministic validation and read-only SQL enforcement.
- MLOps & LLM Delivery: Containerised FastAPI services, Kubernetes/Helm deployment, CI/CD (Jenkins, GitLab, Bitbucket), async persistence with SQLAlchemy and Alembic, structured logging, and cost/quality monitoring in production.
- **Traditional Machine Learning:** Classification, regression, time series analysis, clustering, dimensionality reduction, anomaly detection, mixed-integer optimization, explainability (SHAP, Sobol).
- **Natural Language Processing (NLP):** Sentiment analysis, text classification, topic modelling.
- **Research & Development**
- **Programming and Development Tools:**
  - Python (expert): FastAPI, Pydantic, LangChain, LangGraph, SQLAlchemy, Pandas, NumPy, Scikit-learn, Matplotlib, Kedro, pytest, structlog
  - SQL (advanced), Excel (advanced)
  - Software architecture: hexagonal (ports & adapters), domain-driven design, async services, REST APIs and microservices
- **Cloud Platforms:**
  - Azure: Azure OpenAI, Databricks, Azure SQL, ADF, DevOps
  - AWS: Bedrock, SageMaker, EC2, S3, Athena, EMR, RDS, Redshift
  - GCP: Vertex AI, BigQuery, Cloud Run, Cloud Storage, Cloud SQL, Dataproc, Dataflow, AutoML, Data Fusion
- **Data & Platform Engineering:**
  - PostgreSQL + pgvector, Neo4j, Snowflake, MySQL, PySpark, Apache Airflow, Apache Beam, MLflow
  - Docker, Kubernetes & Helm, Jenkins, GitLab CI, Bitbucket Pipelines, Git, Bash
  - Quality tooling: pytest, ruff, black, mypy, pre-commit, Alembic
- **Experience with agile methodologies:** Azure DevOps, Jira
- **Languages:** Spanish (native). English (advanced). French (limited)

## PROFESSIONAL EXPERIENCE

### AI Engineer | LLM Engineer (2015-Present)

**AI/LLM Engineer with a PhD and 10+ years in applied ML, now focused on production agentic systems: multi-agent architectures, RAG over knowledge graphs and vector stores, and the guardrail, evaluation and observability layers that make them safe enough for regulated enterprise use. Experience across pharmaceuticals, aviation, finance, mining, logistics and retail.**

#### Pharmaceutical & Healthcare

- **Client: Johnson & Johnson (United States of America) via ITMAGINATION (Poland)**

Building agentic AI systems for enterprise architecture governance and self-serve analytics, alongside machine learning for pharmacovigilance.

**ARIA - ARB Intake Verifier:** multi-agent system that automates J&J's Architecture Review Board. It ingests solution-architecture submissions (PPTX, PDF, DOCX and diagrams), reviews them against the enterprise standards library and returns a scored, fully cited report - cutting a review that took a reviewer a full working day down to roughly 10 minutes.

  - **Multi-Agent Architecture:** Four domain architect agents (Business, Technology, Information, Data Science & AI) running on a shared quality-agent engine, orchestrated as a multi-step agentic AI workflow through a LangGraph state machine, with an agent-to-agent (A2A) protocol for cross-domain evidence exchange. Built on LangGraph and the OpenAI Agents SDK.
  - **Multimodal GraphRAG:** Agent reasoning grounded in a Neo4j knowledge graph of the enterprise architecture inventory (LeanIX) and a pgvector store over the standards library, with slide-to-image vision so agents also assess diagrams, and an enforced citation on every finding.

- LLM Safety Framework: Designed and implemented 19 production guardrails at every graph node - prompt injection, PII/PHI redaction (Presidio), data-exfiltration and prompt-leakage detection, schema and citation validation, per-run cost and token ceilings, among others.
- Engineering & Delivery: Async FastAPI service with PostgreSQL, connectors for SharePoint/O365, LeanIX and Atlassian, and Arize Phoenix / OpenTelemetry tracing with per-run cost tracking.
- Tech stack: Python, LangGraph, LangChain, OpenAI Agents SDK, Azure OpenAI (GPT-5.x), FastAPI, Pydantic, Neo4j, PostgreSQL + pgvector, SQLAlchemy, Docker.

**Enterprise Text2SQL Agent:** self-serve natural-language querying of enterprise databases, so business users get answers without depending on a SQL-fluent data analyst.

- Agent Design: A single tool-calling agent that discovers the schema at runtime, writes read-only SQL and recovers from its own errors by receiving database failures back as tool results - onboarding a new database is a connection-string change, with zero schema-specific prompt engineering.
- Hexagonal Architecture & Guardrails: The agent depends only on LLMPort and DatabasePort contracts, so switching LLM vendor or database platform means writing an adapter; read-only SQL validation, step budget and row caps, plus a full offline regression suite driven by a scripted mock LLM (no API spend).
- Tech stack: Python, Azure OpenAI, LangChain, SQLAlchemy, Pydantic, structlog, pytest.

**Pharmacovigilance Risk Models:** machine learning to predict the risks associated with late case reporting to Health Authorities.

- Designed and validated classification models to predict late case reporting, with SHAP and Sobol analysis for explainability and root cause identification, and clustering into confidence buckets with LLM-generated cluster descriptions. Deployed with Data Engineers in Databricks.
- Tech stack: Python, OpenAI, LangChain, Databricks, Bitbucket.

## Aviation & Conversational AI

- **Client: Iberia Airlines (Spain) via Knowmad Mood (Spain)**

Development of **EVA (Enhanced Virtual Assistant)** - a multi-agent conversational ecosystem powered by LLMs to automate Iberia's customer service operations.

- LLM Engineering & Model Optimization: Designed, tuned and deployed OpenAI and AWS Bedrock models for production-grade airline assistants handling check-in, cancellations, flight status, flight documentation and other customer processes.
- Multi-Agent Architecture: Built modular, hexagonal-architecture-based agents (Check-in, Cancellation, Flight Status, Translator, Tone, Booking, Guardrail) using LangChain and FastAPI with structured I/O through Pydantic schemas.
- API Integration: Integrated multiple Iberia APIs and implemented secure, scalable interaction via internal tools and custom MCP servers for external system access.
- Security & Guardrails: Created prompt-injection and data-exfiltration detectors ensuring compliant and safe interactions.
- Evaluation & Monitoring: Developed internal evaluation tools for LLM cost tracking, token efficiency, latency and output quality.
- Tech stack: Python, LangChain, OpenAI, AWS Bedrock (Nova Pro), FastAPI, Pydantic, DynamoDB, structlog, pytest, GitLab CI/CD.

## Geolocation & Mobility Data Science

- **Client: The People Platform (United States of America) via Truelogic Software (Uruguay)**

**TPP Text2SQL:** production conversational analytics platform over US foot-traffic and mobility data, letting business users run their own analyses in natural language instead of queueing for a data analyst who can write SQL.

- Pipeline Architecture: Four-stage hexagonal pipeline (query extraction, context-aware SQL generation, secure execution, natural-language transformation) with session memory and multi-turn follow-up resolution, plus FAISS semantic example selection for few-shot SQL generation.
- Defence-in-Depth Guardrails: A regex input guard, a parallel LLM safety classifier that adds no perceptible latency, and a hardened read-only SQL backstop; validated with a case adversarial battery (injection, jailbreak, exfiltration, obfuscation).
- Quality & Auditability: A case regression suite combining deterministic oracles with an LLM judge, pre-flight validation of extracted entities against the real catalogs, and every turn persisted to a PostgreSQL audit table.
- Earlier work for the same client: scalable algorithms to map US mobility trends and identify home locations, an EMR-based processing pipeline over large geolocation datasets, and the first text2SQL agents (AWS Bedrock Titan, later OpenAI).
- Tech stack: Python, LangChain, OpenAI, FAISS, LangSmith, PostgreSQL, BigQuery, GCS, Streamlit, Docker, Bitbucket Pipelines, PySpark, AWS (EMR, EC2, S3, RDS, Athena, Bedrock), GCP (Dataproc, Dataflow, Cloud Storage, Cloud SQL).

## **Mining**

- **Client: Codelco (Chile) via Option (Chile)**  
Development of a recommendation system to improve exploitation at the Chuquicamata Mine, the world's largest open pit copper mine in terms of excavated volume.
  - Analysed millions of tons of ore to discover trends and patterns, and built KPIs such as extraction uniformity and use of extraction points.
  - Built a mixed linear programming model to optimize the operation of the mine, improving efficiency and output.
  - Built and managed data engineering / data science pipelines with Kedro, leveraging Azure Databricks for scalable processing and Azure Data Factory for orchestration.
  - Tech stack: Python, optimization libraries, Kedro, Azure Databricks, PySpark, MySQL, Git.

## **Retail**

- **Client: Falabella.com (Chile) via Option (Chile)**  
Modeling and transforming warehouse data to support recommendation systems for Falabella.com, the largest retailer in Latin America.
  - Analysed millions of sales transactions, identifying distinct customer clusters and purchasing patterns.
  - Built a collaborative-filtering recommendation system and automated the full pipeline from warehouse extraction to model deployment, with CI/CD.
  - Tech stack: Python, Git, GCP (Vertex AI, Cloud Storage, BigQuery, Cloud Functions, Cloud Composer).

## **Financial & Risk AI Solutions**

- **Client: Pirani Risk (Colombia)**  
Consultant in the development of an Anti-Money Laundering (AML) product, through a two-phased method using PCA and Birch clustering to reduce feature complexity and identify groups within the banks' clients. Tech stack: Python, Git.
- **Client: Energage LLC, TopWorkPlaces.com (United States of America) via Oiga Technologies (Colombia)**  
Consultant in a customer churn prediction project: identified the key features driving unsubscription and built the churn model that informed retention and personalisation decisions. Tech stack: Python, Snowflake, Git.

## **Transport & Logistics**

- **Client: CFL Multimodal (Luxembourg) via University of Luxembourg**  
Development, deployment and automation of ML models to predict disruptions in rail intermodal operations across several EU nations.
  - Built classification and regression models to predict delays in train operations, incorporating mixed integer programming to optimize solutions.
  - NLP Research: Led sentiment analysis projects on customer review data and published two research papers on tourist satisfaction towards transport and urban metro systems, applying LDA topic modelling and classification models.
  - Implemented MLOps to automate data extraction, transformation, model creation and improvement pipelines.
  - Tech stack: Python, optimization libraries, NLTK, Postgres, Docker, Apache Airflow, Git.
- **Clients: MetroValencia and Logitren (Spain) via Technical University of Valencia (Spain)**  
Prediction and optimization of energy consumption in train operation for diesel-electric and electric trains: location-data analysis of Valencia metro and Valencia-Madrid freight services, ML models to predict consumption from driving practices, and optimization of operations to reduce it. Tech stack: Python, Matlab.
- **Clients: Government of Antioquia, Medellin City Hall and Antioquia Railway Department (Colombia)**  
Geospatial data analysis and demand prediction for public infrastructure: railway route feasibility and prioritisation studies, geographical analysis of transport projects, and responsibility for the transport chapter of the Antioquia Development Plan 2020-2023. Tech stack: QGIS, Python.
- **Client: National University of Colombia (Colombia)**  
Data scientist in several transport projects for public agencies in Medellin: analysis of demographic and survey data to understand passenger and freight travel behaviour, and classification models to predict traveller behaviour. Tech stack: R, Python, Matlab, Excel, Biogeme.

---

## EDUCATION

---

|      |                                                                                                  |
|------|--------------------------------------------------------------------------------------------------|
| 2023 | <b>Postdoc – University of Luxembourg</b> – Applied Machine Learning.                            |
| 2019 | <b>Certification in Data Science, Correlation One</b> – Data science for all (DS4A) program.     |
| 2017 | <b>PhD Eng., Technical University of Valencia, Spain</b> – Applied Machine Learning [CUM LAUDE]. |
| 2013 | <b>M.Eng., National University of Colombia</b> – Transportation.                                 |
| 2011 | <b>B.Eng., National University of Colombia</b> – Civil engineering.                              |

---

## PUBLICATIONS

---

[Google Scholar](#)